

CSE 344 Final Review

- Databases
 - Database: Collection of files storing **related data**.
 - Database Management System: an application program that allows us to manage efficiently the collection of data files.
 - Responsibility of DBMS:
 - Data storage and manipulation
 - Black box thought
 - Physical data independence
- Relational Databases
 - Motivation: avoid singular/flat files
 - Data Model: **Mathematical/Conceptual** way for describing the data
 - Schemas and keys
 - Record and attributes
 - Attribute types/typing

Three elements of data models:

- Instance (actual data)
- Schema (what data being stored)
- Query language (how to retrieve and manipulate data)
- Keys: one (or more) attributes that **uniquely identify** a record
 - When multiple keys, we can choose one key as **primary key**.
- Foreign keys: attributes whose value is a key of a record in some other relation.
- SQL Structure:
 - Flat tables:
 - First normal form
 - Crosswalks and joins
 - Breaking up data into multiple relations

Tables are **not ordered**, **flat**, and has **physical data independence** (do not prescribe how they are implemented/stored on disk).

- Code:
 - create statements: declaring **types** and **keys**.
 - insert / delete statements
 - update
 - drop table
- **DISTINCT** : return unique elements
Projection: function which selects a subset of the attributes
- Inner vs. Outer Joining
 - **inner joining**: each row must come from both tables
 - **outer joining**: include rows from only one of the two tables
 - tableA (left/right/full) outer join tableB :
 - Left outer join: includes tuples from tableA even if no match
 - Right outer join: includes tuples from tableB even if no match
 - Full outer join: includes tuples from both tableA and tableB even if no match
- Nested loop semantics
 - Cross-join multiple tables with selection
- Self joins
 - includes one table multiple times
- Aggregation:
 - Everything in SELECT must be either a **GROUP-BY** attribute, or an aggregate.
- WHERE vs. HAVING :
 - WHERE can contain any condition on attributes in FROM clause.
 - Applied to individual rows.
 - No aggregate.
 - HAVING can contain any condition on aggregate expressions and any group-by keys.
 - Entire group is returned, or removed.

- Executing in order of: **FWGHOS**
- Subqueries:
 - In **SELECT** : single attribute projection
 - In **FROM** : use **AS** and **WITH AS**
 - In **WHERE** : use **existential quantifiers** (**EXISTS** , **IN** , **ANY** , **ALL**)
 - **Correlated** subquery: depends on outside query.
 - Finding Witness (the product with max price): compute aggregate in subquery.
- Monotonicity:
 - A query is **monotone** if whenever we add tuples to one or more input tables, the answer to the query will not lose any of the tuples.
 - If Q is a **SELECT-FROM-WHERE** query that **does not** involve any subqueries, Q **is** monotone.
 - Queries with **universal quantifiers**(all) or **negation** must be nested.
- Relational Algebra:
 - Set/Bag semantics: differs on allowing/disallowing duplicates in tuples, also referred as "Standard"/"Extended" Relational Algebra
 - Operators:
 - Union \cup , intersection \cap , difference $-$
 - selection σ
 - projection π
 - Cartesian product \times , join \bowtie
 - Rename ρ
 - Duplicate elimination δ
 - Grouping and aggregation γ
 - sorting τ

All operators take 1 or more relations as input and return another relation.

Join in R.A.:

- **Theta-join**: $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- **Equijoin**: Theta-joins which join condition θ only involves equalities.
- **Natural Join**: $R \bowtie S = \pi_A(\sigma_{\theta}(R \times S))$

- Relational Algebra is **equally expressive** with SQL.

- Datalog:

- For Queries that cannot be defined in RA (recursive queries)
- Terminology:
 - Facts: tuples in database
 - Rules: queries
 - Extensional Predicates: predicates defined in database itself
 - Intensional Predicates: self-defined predicates
 - Head: the intensional predicates to be defined
 - Body: rule for each head (made with atoms/relational predicates)
 - head variables / existential variables
- Fixpoint semantics: start with empty relations, repeat until $IDB_t = IDB_{t+1}$.
 - A datalog program without functions (+, *, ...) always terminates in polynomial time.
- Minimal model semantics: returns the smallest IDB such that all vars in EDB satisfies rule defined for this IDB (recursively too).
- Linear:
 - Right-linear: $T(x, y) :- R(x, z), T(z, y)$
 - Left-linear: $T(x, y) :- T(x, z), R(z, y)$
 - Non-linear: $T(x, y) :- T(x, z), T(z, y)$
- Writing Rules:
 - Safe: if every variable appears in some positive relational atom

- Semi-structured Data

- **Transactional Data:** Frequent read, simple update

Analytical Data: Decision support, multiple joins & group bys

OLTP (online transaction processing): simple look-up, many updates, consistency.

OLAP (online analytical processing): Decision Support

- **Partitioning:** Store data partitioned on multiple servers, easy to write, hard to read (Good for transactional)

Duplication: Duplicate data on multiple files, hard to write, easy to read (Good for analytical)

- **JSON:** semi-structured data model

- `{}` holds objects, `[a,b,...,x]` is an array.

- Duplicate keys are allowed by standard, but not by many implementations.
- Allows null values, duplicate attributes, and nested collections.

- **SQL++:**

- Dataverse: is a database
 - USE + dataverse name
- Type: defines a schema of a collection
 - Lists all required fields, ? for optional fields.
 - OPEN permits other field, while CLOSED doesn't.
 - PRIMARY KEY mykey AUTOGENERATED to autogenerate a unique id for each entry.

- indices:

- Three type of indices available:
 - BTREE : for equality and range queries.
 - RTREE : for two-dimensional range queries:
for example: WHERE x > 20 AND x < 30 AND y > 20 AND y < 30
 - KEYWORD : for substring search.
- **CANNOT** index inside a nested collection.
- Syntax: CREATE INDEX countryID ON country(carCode) TYPE BTREE

- **Heterogeneous collections:** Same attribute, but in some object is array, in some object is not.

- Solution: Use (CASE ... WHEN...THEN...ELSE...) u , and is_array() helper function.

- UNNEST : distribute each element in this array to the rest of the row.

- Multi-value join: split an array to "unnest" with split(y.'country', " ") .

- Physical Plan:

- Overall execution of a SQL query:

Parse SQL -> Select logical plan (Relational Algebra) -> Select physical Plan -> execute.

Physical Plans: Choose an effective implementation for a RA operator.

- Basic Physical Operator: Join

- Nested Loop Join: $O(n^2)$
- Merge Join: $O(n \log(n))$
- Hash Join: $O(n)$ to $O(n^2)$

- Physical Data Independence: Apps are insulated from how data is stored physically.
- Data Storage:
 - DBMS store data in files, each file is splitted into blocks, and each block contains serveral tuples.
 - Data file can be either a Heap file (unsorted), or Sequential file (sorted according to some attribute/key).
 - Index: Additional file mapps from a search key (some attribute's value) to the actual tuple in memory. (*Could have many index for one table.*)
 - Clustered index: records close in index is close in data
 - Unclustered index: records close in index may be far in data
 - Primary: Over attributes including primary key
 - Secondart: Otherwise.
- Scanning a disk: Sequential scan is much faster than random accessing the disk.
- Index should be selected on attributes which:
 - an exact match
 - a range match (clustered index works the best)
 - a join

On such key is needed. (solution to **index selection problem**).

- Cost Estimation:
 - Parameters for cost estimation: for some relation R,
 - $B(R)$: number of blocks in R
 - $T(R)$: number of tuples in R
 - $V(R, a)$: number of distinct values of attribute a
 - Note that when a is a key, $V(R, a) = T(R)$; else, $V(R, a) \leq T(R)$
 - Cost of:
 - Sequentially scan a table: $B(R)$
 - Index-based scanning:
 - Clustered: $f * B(R)$
 - Unclustered: $f * T(R)$
- [Note: f is usually $1/V(R, a)$, which a is the key that index is built on]
- Join Cost:

- **Hash Join:** Scan one table into memory and build hash table, scan in the other table and join.
 - Overall cost: $B(R) + B(S)$
 - Should build hashTable on the relation with more distinct attributes.
 - **Nested Loop Join:** Scan one table into memory, scan in the other and compare each tuple pair to determine if successfully joins.
 - Overall cost: $B(R) + T(R) * B(S)$
 - **Sort-merge Join:** Scan two tables and sort in memory, then merge-join them. (Not one-pass algorithm)
 - Overall cost: $B(R) + B(S)$
 - **Index Nested loop Join:** Assume S has an index page in memory. So iterate over R and fetch corresponding entries of S.
 - Cost if S is clustered: $B(R) + T(R) * (B(S) / V(S, a))$
 - Cost if S is unclustered: $B(R) + T(R) * (T(S) / V(S, a))$
- Parallel/Map-Reduce:
 - Speed up: More nodes(processors, computers), same data, more speed
Scale up: More nodes, more data, same speed
 - Shared-nothing: each machine has its own disk and memory. Easy to maintain and scale, but hard to administer and tune.
 - Intra-operator parallelism:
 - Each operator is placed on multiple nodes.
 - Good for both transactional and analytical work.
 - Horizontal Data Partitioning:
 - Block Partition: partition arbitrarily with block alignment.
 - Hash Partition on some attribute
 - Range Partition: a tuple goes to some server if value falls in some range.
 - Uniform Data/Skewed Data:
 - Uniform Data: Data is partitioned uniformly across servers.
 - Block Partition will be uniform.
 - Skewed Data: Some server holds more data than other servers.
 - Hash Partition/Range Partition on some attribute may cause this problem.
 - Shuffle & Broadcast:

- Shuffle Join: Reshuffle each table on the common attribute, and compute join locally.
- Broadcast join: Reshuffle one table on the common attribute, join with another table, and broadcast the result.
- Map/Reduce framework:
 - Map: extract something you care about for each record.
 - Reduce: aggregate, summarize, filter, transform
 - Fault-tolerance: map writes to local disk, reduce reads from disk. If fails, the reduce task is restarted on another server, but slow.
 - Spark uses "Resilient distributed datasets" = main memory + lineage.
- Straggler: a machine takes unusual long to complete one of the last tasks.
 - Solution: pre-emptive back up execution of last few in-progress tasks.
- Database design:
 - Conceptual model for database design:
 - Entity: an object
 - Attribute: attributes of entity. Each entity set must have a key, which can be one attribute or several attributes.
 - Relationship: how entities are related.
 - Weak entities: entities which key come from other classes which they are related
 - Relationship between entities:
 - **One to one**: No separate table
 - **One to many**: No separate table
 - **Many to many**: Separate table
 - Integrity constraint motivation: a condition specified on a database schema which restricts the data that can be stored on any instance of the database.
 - Key Constraints: the key in each schema must uniquely identifies a person
 - single-value constraints: a person can only have one biological father
 - Referential integrity constraints: if you work for a company, it must exist in the database.
 - Other constraints: value needs to be in legal range.
 - Constraints in SQL :
 - Primary Key, Foreign Key

- `UNIQUE(name)` can be used to ensure `name` is unique (even if it is not primary key). There can be many unique.
 - Foreign key constraints: the referential integrity constraints. The referred column must be a key in another table.
 - Attribute level constraints
 - `NOT NULL`
 - `CHECK` can check any condition
 - Tuple-level constraints
 - Global constraints/Assertions
- Functional Dependencies: if one attribute implies another
- Closure: everything that an attribute determine
- If $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow B, C$.
 - Formal notation: $\{A\}^+ = \{A, B, C\}$
 - If one attribute determines all other attributes in the schema, then it is called **superkey**
 - The smallest subset of attributes that does it is called the **minimal key** or just **key**.

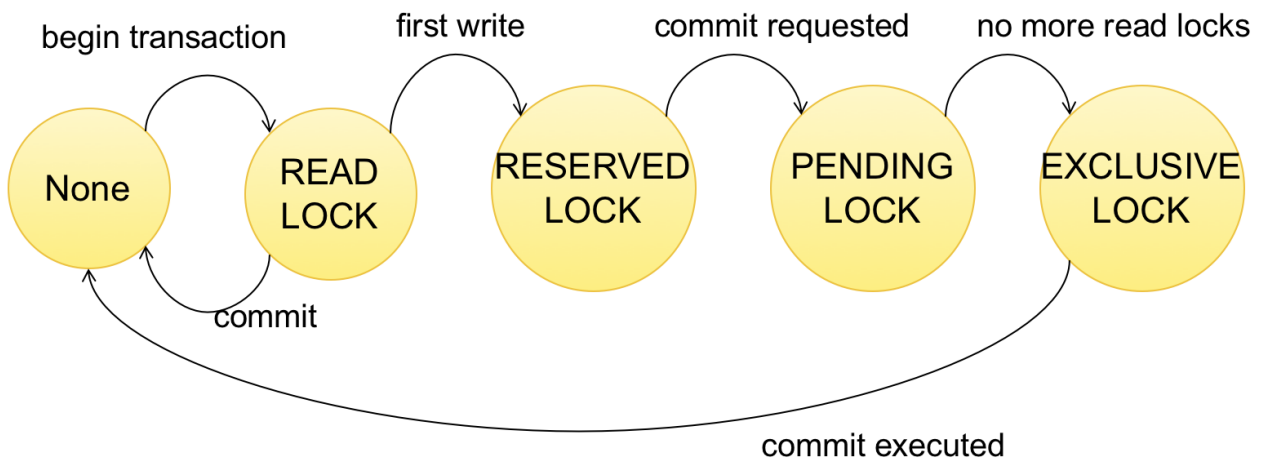
Boyce-Codd Normal Form (BCNF): No bad functional dependencies.

- For all attribute X , either $\{X\}^+ = \{X\}$ or $\{X\}^+ = \text{all attributes}$.
- Whenever $X \rightarrow B$ is a non-trivial dependency, then X is a superkey.

- Transactions:

- Changes to database should be all or nothing.
- Transaction: Collection of statements that are executed atomically (logically speaking)
- Transaction Properties:
 - **Atomic:** State shows either all effect for a transaction, or none of them.
 - **Consistent:** Integrity must hold before and after transaction is executed.
 - **Isolated:** effect of each transaction is as if it were the only transaction running on the system
 - **Durable:** transaction's effect remains after execution finishes.
- **Serial Schedule:** transactions are executed sequentially
- **Serializable Schedule:** equivalent to a serial schedule

- Determine if schedule is serializable by looking at conflicts (swapping will change program behavior)
 - READ-WRITE
 - WRITE-READ
 - WRITE-WRITE
- **Conflict Serializable Schedule:** transform into a serial schedule by swapping adjacent non-conflicting actions. (Every conflict serializable schedule is a serializable schedule).
 - Testing Conflict Serializable Schedule with **Precedence Graph**: each vertex is a transaction; an edge from A to B if an action in A conflicts with action in B and comes before B.
 - The schedule is conflict serializable if and only if the graph is **acyclic**.
- Recoverable: Each transaction commits after all transactions from which it has read has committed.
- **Scheduler** (Concurrency Control Manager): the module which schedules transactions, ensuring serializability.
 - Locking Scheduler: each element acquires the lock, wait until other transactions release the lock, then use. Picture below is the locking scheduler for SQLite.



- If no concurrency control:
 - Dirty Read (including inconsistent read)
 - Unrepeatable reads
 - Lost updates
- Two-phase locking: In every transaction, all lock requests must precede all unlock requests. (Ensures Conflict serializability)

- Strict 2PL:
 - All locks are held until commit/abort:
 - All unlocks are done together with commit/abort.

With strict 2PL, we can get schedules both conflict-serializable and recoverable.

- DEADLOCK: Locks waiting for other locks to release and stuck in infinite loop.
Solution: Use graph to check for cycles and if detected, abort one transaction.
- Isolation Levels in SQL:
 - READ UNCOMMITTED
 - Strict 2PL for WRITE locks, no READ locks.
 - READ COMMITTED
 - Strict 2PL for WRITE locks, only acquire locks for READ locks (not 2PL)
 - REPEATABLE READ
 - Strict 2PL for both READ and WRITE
 - SERIALIZABLE
 - Strict 2PL for both READ and WRITE, and **Predicate locking** to deal with phantoms.
 - Phantom: a tuple that is invisible during part of a transaction execution but not invisible during the entire execution
 - Predicate lock: a lock on an arbitrary predicate.